

Retail Simplilearn Capstone Project 3

March 9, 2023

1 Retail Simplilearn Capstone Project 3

Problem Statement: * It is a critical requirement for business to understand the value derived from a customer. RFM is a method used for analyzing customer value. * Customer segmentation is the practice of segregating the customer base into groups of individuals based on some common characteristics such as age, gender, interests, and spending habits * Perform customer segmentation using RFM analysis. The resulting segments can be ordered from most valuable (highest recency, frequency, and value) to least valuable (lowest recency, frequency, and value).

Dataset Description: This is a transnational data set which contains all the transactions that occurred between 01/12/2010 and 09/12/2011 for a UK-based and registered non-store online retail. The company mainly sells unique and all-occasion gifts.

- **InvoiceNo:** Invoice number. Nominal, a 6-digit integral number uniquely assigned to each transaction. If this code starts with letter 'c', it indicates a cancellation.
- **StockCode:** Product (item) code. Nominal, a 5-digit integral number uniquely assigned to each distinct product.
- **Description:** Product (item) name. Nominal.
- **Quantity:** The quantities of each product (item) per transaction. Numeric.
- **InvoiceDate:** Invoice Date and time. Numeric, the day and time when each transaction was generated.
- **UnitPrice:** Unit price. Numeric, Product price per unit in sterling.
- **CustomerID:** Customer number. Nominal, a 5-digit integral number uniquely assigned to each customer.
- **Country:** Country name. Nominal, the name of the country where each customer resides.

1.0.1 Project Task: Week 1:

Data Cleaning:

1. Perform a preliminary data inspection and data cleaning.
 - a. Check for missing data and formulate an apt strategy to treat them.
 - b. Remove duplicate data records.
 - c. Perform descriptive analytics on the given data.

Data Transformation:

2. Perform cohort analysis (a cohort is a group of subjects that share a defining characteristic). Observe how a cohort behaves across time and compare it to other cohorts.
 - a. Create month cohorts and analyze active customers for each cohort.
 - b. Analyze the retention rate of customers.

1.0.2 Project Task: Week 2

Data Modeling :

1. Build a RFM (Recency Frequency Monetary) model. Recency means the number of days since a customer made the last purchase. Frequency is the number of purchase in a given period. It could be 3 months, 6 months or 1 year. Monetary is the total amount of money a customer spent in that given period. Therefore, big spenders will be differentiated among other customers such as MVP (Minimum Viable Product) or VIP.
2. Calculate RFM metrics.
3. Build RFM Segments. Give recency, frequency, and monetary scores individually by dividing them into quartiles.
 - b1. Combine three ratings to get a RFM segment (as strings).
 - b2. Get the RFM score by adding up the three ratings.
 - b3. Analyze the RFM segments by summarizing them and comment on the findings.

Note: Rate “recency” for customer who has been active more recently higher than the less recent customer, because each company wants its customers to be recent.

Note: Rate “frequency” and “monetary” higher, because the company wants the customer to visit more often and spend more money

1.0.3 Project Task: Week 3

Data Modeling :

1. Create clusters using k-means clustering algorithm.
 - a. Prepare the data for the algorithm. If the data is asymmetrically distributed, manage the skewness with appropriate transformation. Standardize the data.
 - b. Decide the optimum number of clusters to be formed.
 - c. Analyze these clusters and comment on the results.

1.0.4 Project Task: Week 4

Data Reporting:

1. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:

- Country-wise analysis to demonstrate average spend. Use a bar chart to show the monthly figures
- Bar graph of top 15 products which are mostly ordered by the users to show the number of products sold
- Bar graph to show the count of orders vs. hours throughout the day
- Plot the distribution of RFM values using histogram and frequency charts
- Plot error (cost) vs. number of clusters selected
- Visualize to compare the RFM values of the clusters using heatmap

```
>>>>>-----
----->>>>> #
SOLUTION: ## Week 1: ### (A) Data Cleaning (1) Reading Data and Preliminary Data
Inspection
```

```
[5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from datetime import timedelta
from pandas import ExcelWriter
```

```
[6]: df = pd.read_excel("/content/drive/MyDrive/Course 5 - Data Science Capstone_
↳Project/Retail/Project 3/Online Retail.xlsx")
df.head()
```

```
[6]: InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6

InvoiceDate UnitPrice CustomerID Country
0 2010-12-01 08:26:00 2.55 17850.0 United Kingdom
1 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
2 2010-12-01 08:26:00 2.75 17850.0 United Kingdom
3 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
4 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
```

```
[7]: # Check shape of data
df.shape
```

[7]: (541909, 8)

```
[8]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[9]: # Check feature details of data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   InvoiceNo        541909 non-null object
1   StockCode       541909 non-null object
2   Description     540455 non-null object
3   Quantity        541909 non-null int64
4   InvoiceDate     541909 non-null datetime64[ns]
5   UnitPrice       541909 non-null float64
6   CustomerID      406829 non-null float64
7   Country         541909 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
```

- (a) Missing values treatment:

```
[10]: # Check missing values in data
df.isnull().sum()
```

```
[10]: InvoiceNo          0
StockCode            0
Description         1454
Quantity            0
InvoiceDate         0
UnitPrice           0
CustomerID        135080
Country             0
dtype: int64
```

```
[11]: # Calculating the Missing Values % contribution in DF
df_null = round(df.isnull().sum()/len(df)*100,2)
df_null
```

```
[11]: InvoiceNo          0.00
StockCode            0.00
```

```

Description      0.27
Quantity         0.00
InvoiceDate      0.00
UnitPrice        0.00
CustomerID       24.93
Country          0.00
dtype: float64

```

As we can see two columns in data have missing values. * Description - 0.27% (1454 nos.) * CustomerID - 24.93% (135080)

CustomerID is important feature of our analysis since our analysis is centered around Customers only so we can not impute null values **CustomerID** with mean/ median/ mode in this case. We will check possibility to fill null values in **CustomerID** column by looking up for **InvoiceNo** of the row having null **CustomerID** in other rows where **CustomerID** is present. If there are still any null values in **CustomerID** after this process then we will drop complete row having missing **CustomerID**.

We can drop **Description** feature from our data since it is not going to contribute in our model.

```
[12]: invoice_null_custid = set(df[df['CustomerID'].isnull()]['InvoiceNo'])
      df[df['InvoiceNo'].isin(invoice_null_custid) & (~df['CustomerID'].isnull())]
```

```
[12]: Empty DataFrame
      Columns: [InvoiceNo, StockCode, Description, Quantity, InvoiceDate, UnitPrice,
      CustomerID, Country]
      Index: []
```

We could not find any value to impute null values in **CustomerID** column since all entries for a particular **InvoiceNo** have missing **CustomerID** if that particular **InvoiceNo** has null **CustomerID** in even one entry. So we will drop all rows having null values in **CustomerID**.

```
[13]: df = df.drop('Description', axis=1)
      df = df.dropna()
      df.shape
```

```
[13]: (406829, 7)
```

- (b) **Remove duplicate data records:** Since our data is transactional data and it has duplicate entries for InvoiceNo and CustomerID, we will drop only those rows which are completely duplicated, not on the basis of any one particular column such as InvoiceNo or CustomerID etc.

```
[14]: df = df.drop_duplicates()
      df.shape
```

```
[14]: (401602, 7)
```

- (c) **Perform descriptive analysis on the given data:**

```
[15]: # CustomerID is 'float64', changing the datatype of CustomerID to string as
      ↪ Customer ID as numerical data does not make sense

df['CustomerID'] = df['CustomerID'].astype(str)
```

```
[16]: df.describe(datetime_is_numeric=True)
```

```
[16]:
```

	Quantity		InvoiceDate	UnitPrice
count	401602.000000		401602	401602.000000
mean	12.182579	2011-07-10	12:08:08.129839872	3.474064
min	-80995.000000	2010-12-01	08:26:00	0.000000
25%	2.000000	2011-04-06	15:02:00	1.250000
50%	5.000000	2011-07-29	15:40:00	1.950000
75%	12.000000	2011-10-20	11:58:00	3.750000
max	80995.000000	2011-12-09	12:50:00	38970.000000
std	250.283248		NaN	69.764209

- **Quantity:** Average quantity of each product in transaction is 12.18. Also note that minimum value in **Quantity** column is negative. This implies that some customers had returned the product during our period of analysis.
- **InvoiceDate:** Our data has transaction between 01-12-2010 to 09-12-2011
- **UnitPrice:** Average price of each product in transactions is 3.47

```
[17]: df.describe(include=['O'])
```

```
[17]:
```

	InvoiceNo	StockCode	CustomerID	Country
count	401602	401602	401602	401602
unique	22190	3684	4372	37
top	576339	85123A	17841.0	United Kingdom
freq	542	2065	7812	356726

- **InvoiceNo:** Total entries in preprocessed data are 4,01,602 but transactions are 22,190. Most number of entries (count of unique products) are in Invoice No. '576339' and is 542 nos.
- **StockCode:** There are total 3684 unique products in our data and product with stock code '85123A' appears most frequently (2065 times) in our data.
- **CustomerID:** There are 4372 unique customers in our final preprocessed data. Customer with ID '17841' appears most frequently in data (7812 times)
- **Country:** Company has customers across 37 countries. Most entries are from United Kingdom in our dataset (356726)

1.0.5 (B) Data Transformation

(2) Perform Cohort Analysis * (a) Create month cohort of customers and analyze active customers in each cohort:

```
[18]: # Convert to InvoiceDate to Year-Month format
df['month_year'] = df['InvoiceDate'].dt.to_period('M')
```

```
df['month_year'].nunique()
```


[18]: 13

```
[19]: month_cohort = df.groupby('month_year')['CustomerID'].nunique()  
month_cohort
```

```
[19]: month_year  
2010-12    948  
2011-01    783  
2011-02    798  
2011-03   1020  
2011-04    899  
2011-05   1079  
2011-06   1051  
2011-07    993  
2011-08    980  
2011-09   1302  
2011-10   1425  
2011-11   1711  
2011-12    686  
Freq: M, Name: CustomerID, dtype: int64
```

```
[20]: plt.figure(figsize=(10,5))  
sns.barplot(y = month_cohort.index, x = month_cohort.values);  
plt.xlabel("Count of customers")  
plt.title("No. of active customers in each month")
```

[20]: Text(0.5, 1.0, 'No. of active customers in each month')



- (b) Analyze the retention rate of customers:


```
[21]: month_cohort - month_cohort.shift(1)
```

```
[21]: month_year
2010-12      NaN
2011-01    -165.0
2011-02      15.0
2011-03     222.0
2011-04    -121.0
2011-05     180.0
2011-06     -28.0
2011-07    -58.0
2011-08     -13.0
2011-09     322.0
2011-10     123.0
2011-11     286.0
2011-12   -1025.0
Freq: M, Name: CustomerID, dtype: float64
```

```
[22]: retention_rate = round(month_cohort.pct_change( periods=1)*100,2)
      retention_rate
```

```
[22]: month_year
2010-12      NaN
2011-01    -17.41
2011-02      1.92
2011-03     27.82
2011-04    -11.86
2011-05     20.02
2011-06     -2.59
2011-07     -5.52
2011-08     -1.31
2011-09     32.86
2011-10      9.45
2011-11     20.07
2011-12    -59.91
Freq: M, Name: CustomerID, dtype: float64
```

```
[23]: plt.figure(figsize=(10,5))
      sns.barplot(y = retention_rate.index, x = retention_rate.values);
      plt.xlabel("Retention (in %)")
      plt.title("Month-wise customer retention rate");
```

1.1 Week 2:

Monetary analysis:

```
[24]: df['amount'] = df['Quantity']*df['UnitPrice']
df.head()
```

```
[24]: InvoiceNo StockCode Quantity InvoiceDate UnitPrice CustomerID \
0 536365 85123A 6 2010-12-01 08:26:00 2.55 17850.0
1 536365 71053 6 2010-12-01 08:26:00 3.39 17850.0
2 536365 84406B 8 2010-12-01 08:26:00 2.75 17850.0
3 536365 84029G 6 2010-12-01 08:26:00 3.39 17850.0
4 536365 84029E 6 2010-12-01 08:26:00 3.39 17850.0
```

```
Country month_year amount
0 United Kingdom 2010-12 15.30
1 United Kingdom 2010-12 20.34
2 United Kingdom 2010-12 22.00
3 United Kingdom 2010-12 20.34
4 United Kingdom 2010-12 20.34
```

```
[25]: df_monetary = df.groupby('CustomerID').sum()['amount'].reset_index()
df_monetary
```

```
[25]: CustomerID amount
0 12346.0 0.00
1 12347.0 4310.00
```

2	12348.0	1797.24
3	12349.0	1757.55
4	12350.0	334.40
...
4367	18280.0	180.60
4368	18281.0	80.82
4369	18282.0	176.60
4370	18283.0	2045.53
4371	18287.0	1837.28

[4372 rows x 2 columns]

Frequency Analysis:

```
[26]: df_frequency = df.groupby('CustomerID').nunique()['InvoiceNo'].reset_index()
# df_frequency = df.drop_duplicates('InvoiceNo').groupby('CustomerID').
# count()['InvoiceNo'].reset_index()
df_frequency
```

```
[26]:      CustomerID  InvoiceNo
0         12346.0          2
1         12347.0          7
2         12348.0          4
3         12349.0          1
4         12350.0          1
...          ...          ...
4367        18280.0          1
4368        18281.0          1
4369        18282.0          3
4370        18283.0         16
4371        18287.0          3
```

[4372 rows x 2 columns]

Recency Analysis:

```
[27]: # We will fix reference date for calculating recency as last transaction day in
# data + 1 day
ref_day = max(df['InvoiceDate']) + timedelta(days=1)
df['days_to_last_order'] = (ref_day - df['InvoiceDate']).dt.days
df.head()
```

```
[27]:      InvoiceNo  StockCode  Quantity      InvoiceDate  UnitPrice  CustomerID  \
0      536365      85123A         6 2010-12-01 08:26:00         2.55      17850.0
1      536365       71053         6 2010-12-01 08:26:00         3.39      17850.0
2      536365      84406B         8 2010-12-01 08:26:00         2.75      17850.0
3      536365      84029G         6 2010-12-01 08:26:00         3.39      17850.0
4      536365      84029E         6 2010-12-01 08:26:00         3.39      17850.0
```

	Country	month_year	amount	days_to_last_order
0	United Kingdom	2010-12	15.30	374
1	United Kingdom	2010-12	20.34	374
2	United Kingdom	2010-12	22.00	374
3	United Kingdom	2010-12	20.34	374
4	United Kingdom	2010-12	20.34	374

```
[28]: df_recency = df.groupby('CustomerID')['days_to_last_order'].min().reset_index()
df_recency
```

```
[28]:
```

	CustomerID	days_to_last_order
0	12346.0	326
1	12347.0	2
2	12348.0	75
3	12349.0	19
4	12350.0	310
...
4367	18280.0	278
4368	18281.0	181
4369	18282.0	8
4370	18283.0	4
4371	18287.0	43

[4372 rows x 2 columns]

Calculate RFM metrics:

```
[29]: df_rf = pd.merge(df_recency, df_frequency, on='CustomerID', how='inner')
df_rfm = pd.merge(df_rf, df_monetary, on='CustomerID', how='inner')
df_rfm.columns = ['CustomerID', 'Recency', 'Frequency', 'Monetary']
df_rfm.head()
```


```
[29]:
```

	CustomerID	Recency	Frequency	Monetary
0	12346.0	326	2	0.00
1	12347.0	2	7	4310.00
2	12348.0	75	4	1797.24
3	12349.0	19	1	1757.55
4	12350.0	310	1	334.40

Build RFM Segments:


```
[30]: df_rfm['recency_labels'] = pd.cut(df_rfm['Recency'], bins=5,
                                         labels=['newest', 'newer', 'medium', 'old',
                                         ↪ 'older', 'oldest'])
df_rfm['recency_labels'].value_counts().plot(kind='barh');
df_rfm['recency_labels'].value_counts()
```

```
[30]: newest      2734
      newer       588
      medium     416
      older      353
      oldest     281
      Name: recency_labels, dtype: int64
```




```
[31]: df_rfm['frequency_labels'] = pd.cut(df_rfm['Frequency'], bins=5,
      ↪ labels=['lowest', 'lower', 'medium', 'higher', 'highest'])
      df_rfm['frequency_labels'].value_counts().plot(kind='barh');
      df_rfm['frequency_labels'].value_counts()
```

```
[31]: lowest      4348
      lower       18
      medium       3
      highest      2
      higher       1
      Name: frequency_labels, dtype: int64
```



```
[32]: df_rfm['monetary_labels'] = pd.cut(df_rfm['Monetary'], bins=5,
    ↪ labels=['smallest', 'smaller', 'medium', 'larger', 'largest'])
df_rfm['monetary_labels'].value_counts().plot(kind='barh');
df_rfm['monetary_labels'].value_counts()
```

```
[32]: smallest    4357
smaller         9
medium          3
largest         2
larger          1
Name: monetary_labels, dtype: int64
```



```
[33]: df_rfm['rfm_segment'] =
    ↳ df_rfm[['recency_labels', 'frequency_labels', 'monetary_labels']].agg('-',
    ↳ join, axis=1)
df_rfm.head()
```

```
[33]: CustomerID  Recency  Frequency  Monetary  recency_labels  frequency_labels \
0    12346.0      326         2         0.00      oldest      lowest
1    12347.0         2         7    4310.00      newest      lowest
2    12348.0        75         4    1797.24      newest      lowest
3    12349.0        19         1    1757.55      newest      lowest
4    12350.0       310         1     334.40      oldest      lowest
```

```
monetary_labels      rfm_segment
0      smallest  oldest-lowest-smallest
1      smallest  newest-lowest-smallest
2      smallest  newest-lowest-smallest
3      smallest  newest-lowest-smallest
4      smallest  oldest-lowest-smallest
```

RFM Score:

```
[34]: recency_dict = {'newest': 5, 'newer':4, 'medium': 3, 'older':2, 'oldest':1}
frequency_dict = {'lowest':1, 'lower':2, 'medium': 3, 'higher':4, 'highest':5}
monetary_dict = {'smallest':1, 'smaller':2, 'medium': 3, 'larger':4, 'largest':
    ↳ 5}
```


```
df_rfm['rfm_score'] = df_rfm['recency_labels'].map(recency_dict).astype(int)+  
↳df_rfm['frequency_labels'].map(frequency_dict).astype(int) +  
↳df_rfm['monetary_labels'].map(monetary_dict).astype(int)  
df_rfm.head(10)
```

```
[34]: CustomerID Recency Frequency Monetary recency_labels frequency_labels \  
0    12346.0      326         2         0.00      oldest      lowest  
1    12347.0         2         7    4310.00      newest      lowest  
2    12348.0        75         4    1797.24      newest      lowest  
3    12349.0        19         1    1757.55      newest      lowest  
4    12350.0       310         1     334.40      oldest      lowest  
5    12352.0        36        11    1545.41      newest      lowest  
6    12353.0       204         1      89.00      medium      lowest  
7    12354.0       232         1    1079.40      older       lowest  
8    12355.0       214         1     459.40      medium      lowest  
9    12356.0        23         3    2811.43      newest      lowest
```


```
monetary_labels      rfm_segment  rfm_score  
0      smallest  oldest-lowest-smallest      3  
1      smallest  newest-lowest-smallest      7  
2      smallest  newest-lowest-smallest      7  
3      smallest  newest-lowest-smallest      7  
4      smallest  oldest-lowest-smallest      3  
5      smallest  newest-lowest-smallest      7  
6      smallest  medium-lowest-smallest      5  
7      smallest  older-lowest-smallest      4  
8      smallest  medium-lowest-smallest      5  
9      smallest  newest-lowest-smallest      7
```

Analyze RFM Segment and Score:

```
[35]: df_rfm['rfm_segment'].value_counts().plot(kind='barh', figsize=(10, 5));
```



```
[36]: df_rfm['rfm_score'].value_counts().plot(kind='barh', figsize=(10, 5));
```



1.2 Week 3

1.2.1 Data Modeling:

1. Create clusters using k-means clustering algorithm.

a. Prepare the data for the algorithm. If the data is asymmetrically distributed, manage the skewness with appropriate transformation. Standardize the data.

```
[37]: print(df_rfm.shape)
df_rfm.head()
```

(4372, 9)

```
[37]: CustomerID  Recency  Frequency  Monetary  recency_labels  frequency_labels  \
0    12346.0      326         2         0.00         oldest          lowest
1    12347.0         2         7    4310.00         newest          lowest
2    12348.0        75         4    1797.24         newest          lowest
3    12349.0        19         1    1757.55         newest          lowest
4    12350.0       310         1        334.40         oldest          lowest

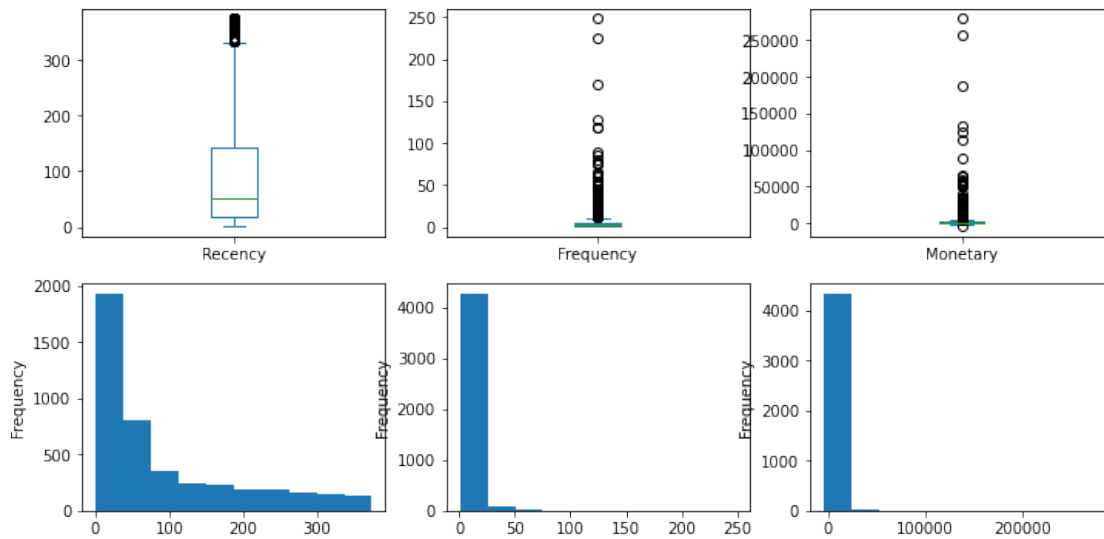
monetary_labels      rfm_segment  rfm_score
0      smallest  oldest-lowest-smallest      3
1      smallest  newest-lowest-smallest      7
2      smallest  newest-lowest-smallest      7
```



```
3         smallest  newest-lowest-smallest      7
4         smallest  oldest-lowest-smallest     3
```

```
[38]: plt.figure(figsize=(12,6))

for i, feature in enumerate(['Recency', 'Frequency', 'Monetary']):
    plt.subplot(2,3,i+1)
    df_rfm[feature].plot(kind='box')
    plt.subplot(2,3,i+1+3)
    df_rfm[feature].plot(kind='hist')
```



Outliers: Frequency and Monetary features in above data seem to have lot of outliers. Lets drop them.

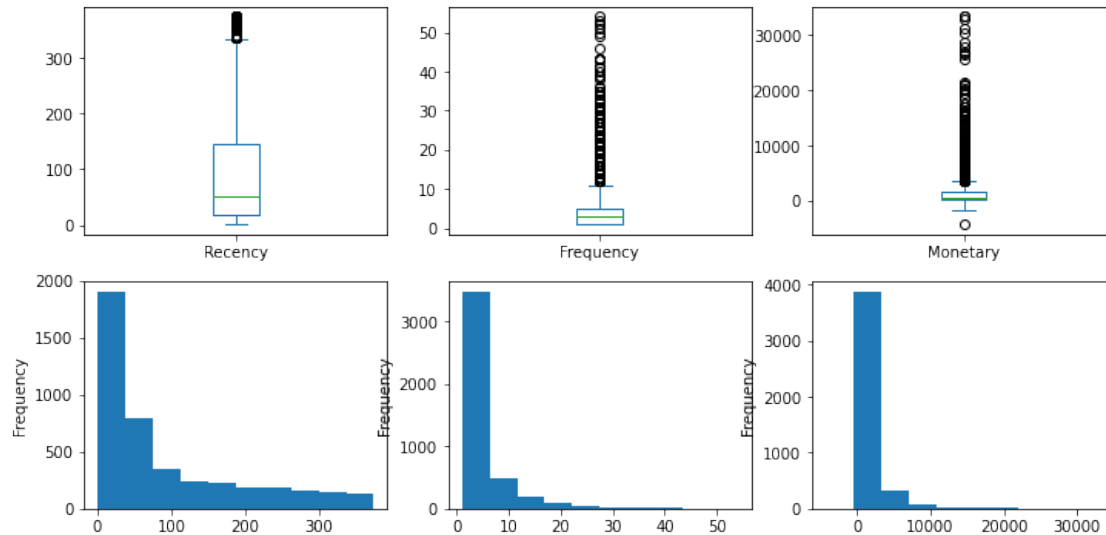
```
[39]: df_rfm = df_rfm[(df_rfm['Frequency']<60) & (df_rfm['Monetary']<40000)]
df_rfm.shape
```

```
[39]: (4346, 9)
```

26 Customers removed as outlier from out data.

```
[40]: plt.figure(figsize=(12,6))

for i, feature in enumerate(['Recency', 'Frequency', 'Monetary']):
    plt.subplot(2,3,i+1)
    df_rfm[feature].plot(kind='box')
    plt.subplot(2,3,i+1+3)
    df_rfm[feature].plot(kind='hist')
```



Log Transformation: Now since all three features have right skewed data therefore we will use log transformation of these features in our model.

```
[41]: df_rfm_log_trans = pd.DataFrame()
df_rfm_log_trans['Recency'] = np.log(df_rfm['Recency'])
df_rfm_log_trans['Frequency'] = np.log(df_rfm['Frequency'])
df_rfm_log_trans['Monetary'] = np.log(df_rfm['Monetary']-df_rfm['Monetary'].
    ↳min()+1)
```

Standard Scalar Transformation: It is extremely important to rescale the features so that they have a comparable scale.

```
[42]: scaler = StandardScaler()

df_rfm_scaled = scaler.fit_transform(df_rfm_log_trans[['Recency', 'Frequency', 'Monetary']])
df_rfm_scaled

df_rfm_scaled = pd.DataFrame(df_rfm_scaled)
df_rfm_scaled.columns = ['Recency', 'Frequency', 'Monetary']
df_rfm_scaled.head()
```

```
[42]:   Recency  Frequency  Monetary
0   1.402988  -0.388507  -0.770922
1  -2.100874   0.967301   1.485132
2   0.392218   0.361655   0.364190
3  -0.552268  -1.138669   0.342970
4   1.368370  -1.138669  -0.527416
```

b. Build K-Means Clustering Model and Decide the optimum number of clusters to

be formed.

```
[43]: # k-means with some arbitrary k
kmeans = KMeans(n_clusters=3, max_iter=50)
kmeans.fit(df_rfm_scaled)
```

```
[43]: KMeans(max_iter=50, n_clusters=3)
```


```
[44]: kmeans.labels_
```

```
[44]: array([0, 1, 2, ..., 2, 1, 2], dtype=int32)
```

```
[45]: # Finding the Optimal Number of Clusters with the help of Elbow Curve/ SSD
ssd = []
range_n_clusters = [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]
for num_clusters in range_n_clusters:
    kmeans = KMeans(n_clusters=num_clusters, max_iter=100)
    kmeans.fit(df_rfm_scaled)

    ssd.append(kmeans.inertia_)

# plot the SSDs for each n_clusters
plt.plot(range_n_clusters,ssd);
```



```
[46]: # Creating dataframe for exporting to create visualization in tableau later
```

```
df_inertia = pd.DataFrame(list(zip(range_n_clusters, ssd)),  
    columns=['clusters', 'intertia'])  
df_inertia
```

```
[46]:
```

	clusters	intertia
0	2	7113.097396
1	3	5342.830291
2	4	4481.006767
3	5	3730.673148
4	6	3044.765927
5	7	2598.358198
6	8	2299.101361
7	9	2044.744830
8	10	1852.941597
9	11	1700.371054
10	12	1575.500092

```
[47]: # Finding the Optimal Number of Clusters with the help of Silhouette Analysis  
range_n_clusters = [2, 3, 4, 5, 6, 7, 8, 9, 10]  
  
for num_clusters in range_n_clusters:  
    kmeans = KMeans(n_clusters=num_clusters, max_iter=50)  
    kmeans.fit(df_rfm_scaled)  
  
    cluster_labels = kmeans.labels_  
  
    silhouette_avg = silhouette_score(df_rfm_scaled, cluster_labels)  
    print("For n_clusters={0}, the silhouette score is {1}".  
        format(num_clusters, silhouette_avg))
```

```
For n_clusters=2, the silhouette score is 0.44132753537785846  
For n_clusters=3, the silhouette score is 0.38147013703040916  
For n_clusters=4, the silhouette score is 0.36228025157331784  
For n_clusters=5, the silhouette score is 0.36790221137230344  
For n_clusters=6, the silhouette score is 0.3523217564861644  
For n_clusters=7, the silhouette score is 0.34283740266136054  
For n_clusters=8, the silhouette score is 0.33524563896109344  
For n_clusters=9, the silhouette score is 0.346301798458803  
For n_clusters=10, the silhouette score is 0.3561080611933406
```

We can select optimum number of clusters as 3 in our final model

```
[48]: # Final model with k=3  
kmeans = KMeans(n_clusters=3, max_iter=50)  
kmeans.fit(df_rfm_scaled)
```

```
[48]: KMeans(max_iter=50, n_clusters=3)
```


c. Analyze these clusters and comment on the results.

```
[49]: # assign the label
df_rfm['Cluster_Id'] = kmeans.labels_
df_rfm.head()
```


```
[49]:  CustomerID  Recency  Frequency  Monetary  recency_labels  frequency_labels  \
0    12346.0     326         2         0.00         oldest         lowest
1    12347.0         2         7    4310.00         newest         lowest
2    12348.0        75         4    1797.24         newest         lowest
3    12349.0        19         1    1757.55         newest         lowest
4    12350.0     310         1     334.40         oldest         lowest
```

	monetary_labels	rfm_segment	rfm_score	Cluster_Id
0	smallest	oldest-lowest-smallest	3	0
1	smallest	newest-lowest-smallest	7	1
2	smallest	newest-lowest-smallest	7	2
3	smallest	newest-lowest-smallest	7	0
4	smallest	oldest-lowest-smallest	3	0


```
[50]: # Box plot to visualize Cluster Id vs Monetary
sns.boxplot(x='Cluster_Id', y='Monetary', data=df_rfm);
```



```
[51]: # Box plot to visualize Cluster Id vs Frequency
sns.boxplot(x='Cluster_Id', y='Frequency', data=df_rfm);
```



```
[52]: # Box plot to visualize Cluster Id vs Recency
sns.boxplot(x='Cluster_Id', y='Recency', data=df_rfm);
```



1.2.2 Inference:

As we can observe from above boxplots that our model has nicely created 3 segments of customer with the interpretation as below: * Customers with Cluster Id 0 are less frequent buyers with low monetary expenditure and also they have not purchased anything in recent time and hence least important for business. * Customers with Cluster Id 1 are the customers having Recency, Frequency and Monetary score in the medium range. * Customers with Cluster Id 2 are the most frequent buyers, spending high amount and recently placing orders so they are the most important customers from business point of view.

1.3 Week 4:

Data Reporting: 1. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:

- a. Country-wise analysis to demonstrate average spend. Use a bar chart to show the monthly
- b. Bar graph of top 15 products which are mostly ordered by the users to show the number of
- c. Bar graph to show the count of orders vs. hours throughout the day
- d. Plot the distribution of RFM values using histogram and frequency charts
- e. Plot error (cost) vs. number of clusters selected
- f. Visualize to compare the RFM values of the clusters using heatmap

```
[54]: !pip install xlswriter
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-  
wheels/public/simple/
```

```
Collecting xlswriter
```

```
  Downloading XlsxWriter-3.0.7-py3-none-any.whl (152 kB)
```

```
152.8/152.8 KB
```

```
13.4 MB/s eta 0:00:00
```

```
Installing collected packages: xlswriter
```

```
Successfully installed xlswriter-3.0.7
```

```
[55]: # Writing dataframe to excel file for creating visualization in tableau  
writer = pd.ExcelWriter('C:\\Users\\mgupt\\mgpython\\Capstone Project\\Retail -  
→PGP\\output_data.xlsx', engine='xlswriter')  
  
df.to_excel(writer, sheet_name='master_data', index=False)  
df_rfm.to_excel(writer, sheet_name='rfm_data', index=False)  
df_inertia.to_excel(writer, sheet_name='inertia', index=False)  
writer.save()
```

```
[58]: product_desc = pd.read_excel("/content/drive/MyDrive/Course 5 - Data Science_  
→Capstne Project/Retail/Project 3/Online Retail.xlsx")  
product_desc = product_desc[['StockCode', 'Description']]  
product_desc = product_desc.drop_duplicates()  
product_desc.to_csv('product_desc.csv', index=False)
```

1.4 *Please refer Dashboard created in Tableau for visualization and graphs*

Link : https://public.tableau.com/app/profile/santhosh.tn/viz/DashboardforSimplilearnCapstoneProject-3Retail_16783246338670/Dashboard1?publish=yes

[]: